

IMPORTING LIBRARIES

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [ ]: import os
os.chdir("C:/Datasets")
```

```
In [ ]: df = pd.read_csv('abalone.csv')
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.23883
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.13920
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.00150
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.13000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.23400
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.32900
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.00500

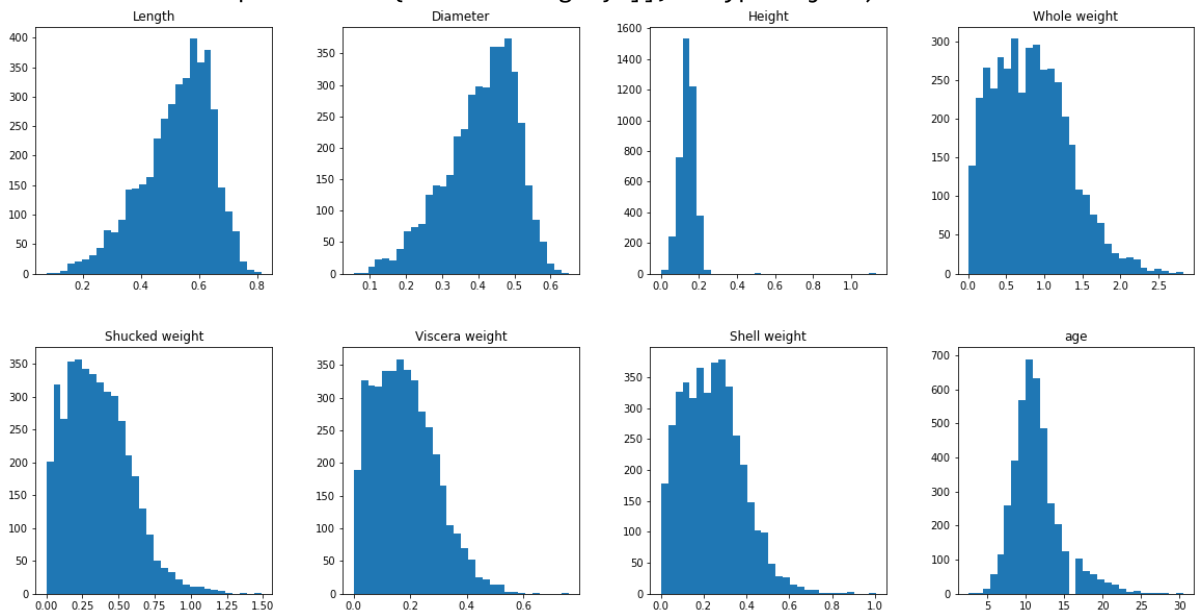
```
In [ ]: df['age'] = df['Rings'] + 1.5
df = df.drop('Rings', axis = 1)
```

```
In [ ]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
from sklearn.metrics import r2_score, mean_squared_error
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

UNIVARIATE ANALYSIS


```
In [ ]: df.hist(figsize=(20,10), grid=False, layout=(2, 4), bins = 30)
```

```
Out[ ]: array([[<AxesSubplot:title={'center':'Length'}>,
      <AxesSubplot:title={'center':'Diameter'}>,
      <AxesSubplot:title={'center':'Height'}>,
      <AxesSubplot:title={'center':'Whole weight'}>],
      [<AxesSubplot:title={'center':'Shucked weight'}>,
      <AxesSubplot:title={'center':'Viscera weight'}>,
      <AxesSubplot:title={'center':'Shell weight'}>,
      <AxesSubplot:title={'center':'age'}>]], dtype=object)
```



```
In [ ]: sns.heatmap(df.isnull())
```

```
Out[ ]: <AxesSubplot:>
```




```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Sex                    4177 non-null   object
1   Length                 4177 non-null   float64
2   Diameter               4177 non-null   float64
3   Height                 4177 non-null   float64
4   Whole weight           4177 non-null   float64
5   Shucked weight         4177 non-null   float64
6   Viscera weight         4177 non-null   float64
7   Shell weight           4177 non-null   float64
8   age                    4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```

```
In [ ]: sns.countplot(x = 'Sex', data = df, palette = 'Set3')
```

```
Out[ ]: <AxesSubplot:xlabel='Sex', ylabel='count'>
```



```
In [ ]: plt.figure(figsize = (20,7))
sns.swarmplot(x = 'Sex', y = 'age', data = df, hue = 'Sex')
sns.violinplot(x = 'Sex', y = 'age', data = df)
```

C:\Users\Harini S\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 56.2% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)


C:\Users\Harini S\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 52.2% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

C:\Users\Harini S\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 58.5% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.


warnings.warn(msg, UserWarning)

```
Out[ ]: <AxesSubplot:xlabel='Sex', ylabel='age'>
```



```
In [ ]: sns.histplot(df.age,kde=True)
```

```
Out[ ]: <AxesSubplot:xlabel='age', ylabel='Count'>
```



BIVARIATE ANALYSIS

```
In [ ]: num_fea = df.select_dtypes(include = [np.number]).columns
ctg_fea = df.select_dtypes(include = [np.object]).columns
```

```
In [ ]: num_fea
```

```
Out[ ]: Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
            'Viscera weight', 'Shell weight', 'age'],
            dtype='object')
```

```
In [ ]: plt.figure(figsize = (20,7))
sns.heatmap(df[num_fea].corr(),annot = True)
```


```
Out[ ]: <AxesSubplot:>
```



MULTI VARIATE ANALYSIS

```
In [ ]: sns.pairplot(df)
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x28913a93fa0>
```



MISSING VALUES

```
In [ ]: missing_values = df.isnull().sum()
```

```
In [ ]: missing_values
```

```
Out[ ]: Sex                0
Length                0
Diameter              0
Height                0
Whole weight          0
Shucked weight        0
Viscera weight        0
Shell weight          0
age                   0
dtype: int64
```

```
In [ ]: missing_values = df.isnull().sum().sort_values(ascending = False)
percentage_missing_values = (missing_values/len(df))*100
pd.concat([missing_values, percentage_missing_values], axis = 1, keys= ['Missing v
```

```
Out[ ]: 
```


	Missing values	% Missing
Sex	0	0.0
Length	0	0.0
Diameter	0	0.0
Height	0	0.0
Whole weight	0	0.0
Shucked weight	0	0.0
Viscera weight	0	0.0
Shell weight	0	0.0
age	0	0.0

OUTLIERS

```
In [ ]: df = pd.get_dummies(df)
dummy_data = df.copy()
```


```
In [ ]: df.boxplot( rot = 90, figsize=(20,5))
```

```
Out[ ]: <AxesSubplot:>
```




```
In [ ]: var = 'Viscera weight'
plt.scatter(x = df[var], y = df['age'],)
```

```
plt.grid(True)
```




```
In [ ]: # outliers removal
df.drop(df[(df['Viscera weight'] > 0.5) & (df['age'] < 20)].index, inplace=True)
df.drop(df[(df['Viscera weight'] < 0.5) & (df['age'] > 25)].index, inplace=True)
```

```
In [ ]: var = 'Shell weight'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```




```
In [ ]: df.drop(df[(df['Shell weight'] > 0.6) & (df['age'] < 25)].index, inplace=True)
df.drop(df[(df['Shell weight'] < 0.8) & (df['age'] > 25)].index, inplace=True)
```

```
In [ ]: var = 'Shucked weight'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```




```
In [ ]: df.drop(df[(df['Shucked weight'] >= 1) & (df['age'] < 20)].index, inplace=True)
df.drop(df[(df['Shucked weight'] < 1) & (df['age'] > 20)].index, inplace=True)
```

```
In [ ]: var = 'Whole weight'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```




```
In [ ]: df.drop(df[(df['Whole weight'] >= 2.5) & (df['age'] < 25)].index, inplace=True)
df.drop(df[(df['Whole weight'] < 2.5) & (df['age'] > 25)].index, inplace=True)
```

```
In [ ]: var = 'Diameter'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```





```
In [ ]: df.drop(df[(df['Diameter']<0.1) & (df['age'] < 5)].index, inplace=True)
df.drop(df[(df['Diameter']<0.6) & (df['age'] > 25)].index, inplace=True)
df.drop(df[(df['Diameter']>=0.6) & (df['age']< 25)].index, inplace=True)
```

```
In [ ]: var = 'Height'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```



```
In [ ]: df.drop(df[(df['Height']>0.4) & (df['age'] < 15)].index, inplace=True)
df.drop(df[(df['Height']<0.4) & (df['age'] > 25)].index, inplace=True)
```

```
In [ ]: var = 'Length'
plt.scatter(x = df[var], y = df['age'],)
plt.grid(True)
```



```
In [ ]: df.drop(df[(df['Length']<0.1) & (df['age'] < 5)].index, inplace=True)
df.drop(df[(df['Length']<0.8) & (df['age'] > 25)].index, inplace=True)
df.drop(df[(df['Length']>=0.8) & (df['age']< 25)].index, inplace=True)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3995 entries, 0 to 4176
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Length                 3995 non-null   float64
1   Diameter               3995 non-null   float64
2   Height                 3995 non-null   float64
3   Whole weight           3995 non-null   float64
4   Shucked weight         3995 non-null   float64
5   Viscera weight         3995 non-null   float64
6   Shell weight           3995 non-null   float64
7   age                    3995 non-null   float64
8   Sex_F                  3995 non-null   uint8
9   Sex_I                  3995 non-null   uint8
10  Sex_M                  3995 non-null   uint8
dtypes: float64(8), uint8(3)
memory usage: 292.6 KB
```

In []: df

Out[]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	16.5	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	8.5	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	10.5	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	11.5	0	0	1
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	8.5	0	1	0
...
4172	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	12.5	1	0	0
4173	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	11.5	0	0	1
4174	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	10.5	0	0	1
4175	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	11.5	1	0	0
4176	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	13.5	0	0	1

3995 rows × 11 columns

CATEGORICAL COLUMNS

In []: cat

Out[]: pandas.core.arrays.categorical.Categorical

```
In [ ]: num_fea = df.select_dtypes(include = [np.number]).columns
ctg_fea = df.select_dtypes(include = [np.object]).columns
```

In []: num_fea

Out[]: Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight', 'Viscera weight', 'Shell weight', 'age', 'Sex_F', 'Sex_I', 'Sex_M'], dtype='object')

```
In [ ]: df_numeric = df[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
```

```
In [ ]: df_numeric.head()
```

```
Out[ ]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5	0	0	1
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5	0	1	0

```
In [ ]: ctg_fea
```

```
Out[ ]: Index([], dtype='object')
```

INDEPENDENT AND DEPENDENT VARIABLE

```
In [ ]: x = df.iloc[:, 0:1].values
```

```
In [ ]: y = df.iloc[:, 1]
```

```
In [ ]: x
```

```
Out[ ]: array([[0.455],
               [0.35 ],
               [0.53 ],
               ...,
               [0.6   ],
               [0.625],
               [0.71 ]])
```

```
In [ ]: y
```

```
Out[ ]: 0      0.365
1      0.265
2      0.420
3      0.365
4      0.255
...
4172   0.450
4173   0.440
4174   0.475
4175   0.485
4176   0.555
Name: Diameter, Length: 3995, dtype: float64
```

SCALING THE INDEPENDENT VARIABLE

```
In [ ]: print ("\n ORIGINAL VALUES: \n\n", x,y)
```

ORIGINAL VALUES:

```
[[0.455]
 [0.35 ]
 [0.53 ]
 ...
 [0.6 ]
 [0.625]
 [0.71 ]] 0      0.365
1      0.265
2      0.420
3      0.365
4      0.255
...
4172   0.450
4173   0.440
4174   0.475
4175   0.485
4176   0.555
Name: Diameter, Length: 3995, dtype: float64
```

```
In [ ]: from sklearn import preprocessing
min_max_scaler = preprocessing.MinMaxScaler(feature_range =(0, 1))
new_y= min_max_scaler.fit_transform(x,y)
print ("\n VALUES AFTER MIN MAX SCALING: \n\n", new_y)
```

VALUES AFTER MIN MAX SCALING:

```
[[0.51587302]
 [0.34920635]
 [0.63492063]
 ...
 [0.74603175]
 [0.78571429]
 [0.92063492]]
```

SPLITTING THE DATA

```
In [ ]: X = df.drop('age', axis = 1)
y = df['age']
```

```
In [ ]: standardScale = StandardScaler()
standardScale.fit_transform(X)

selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)

X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)
```

BUILDING MODEL

LINEAR REGRESSION

```
In [ ]: from sklearn.linear_model import LinearRegression
```

```
In [ ]: lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
Out[ ]: LinearRegression()
```

```
In [ ]: y_train_pred = lm.predict(X_train)
        y_test_pred = lm.predict(X_test)
```

TRAINING THE MODEL

```
In [ ]: X_train
```

```
Out[ ]: array([[0.35 , 0.265, 0.095, ..., 0.    , 1.    , 0.    ],
               [0.465, 0.37 , 0.12 , ..., 0.    , 1.    , 0.    ],
               [0.435, 0.335, 0.1  , ..., 0.    , 1.    , 0.    ],
               ...,
               [0.515, 0.395, 0.125, ..., 0.    , 1.    , 0.    ],
               [0.515, 0.38 , 0.12 , ..., 0.    , 1.    , 0.    ],
               [0.37 , 0.275, 0.1  , ..., 0.    , 1.    , 0.    ]])
```

```
In [ ]: y_train
```

```
Out[ ]: 3431    6.5
        1566   10.5
        1559    8.5
        1284   10.5
         41   15.5
        ...
       2896    9.5
         2   10.5
       1465    9.5
       1290    8.5
       3107    6.5
Name: age, Length: 2996, dtype: float64
```

```
In [ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error
        s = mean_squared_error(y_train, y_train_pred)
        print('Mean Squared error of training set :%2f'%s)
```

```
Mean Squared error of training set :3.499447
```

TESTING THE MODEL

```
In [ ]: X_test
```

```
Out[ ]: array([[0.575, 0.45 , 0.16 , ..., 1.    , 0.    , 0.    ],
               [0.255, 0.195, 0.07 , ..., 0.    , 1.    , 0.    ],
               [0.41 , 0.33 , 0.105, ..., 0.    , 1.    , 0.    ],
               ...,
               [0.415, 0.325, 0.105, ..., 1.    , 0.    , 0.    ],
               [0.495, 0.385, 0.125, ..., 0.    , 1.    , 0.    ],
               [0.36 , 0.265, 0.075, ..., 0.    , 1.    , 0.    ]])
```

```
In [ ]: y_test
```

```
Out[ ]: 1137    11.5
        464     7.5
        3885    8.5
        80     10.5
        3213   14.5
        ...
        248     8.5
        813     6.5
        212    13.5
        2299    9.5
        3529    7.5
Name: age, Length: 999, dtype: float64
```

```
In [ ]: p = mean_squared_error(y_test, y_test_pred)
        print('Mean Squared error of testing set :%2f'%p)
```

Mean Squared error of testing set :3.738376

```
In [ ]: from sklearn.metrics import r2_score
        s = r2_score(y_train, y_train_pred)
        print('R2 Score of training set:%.2f'%s)
```

R2 Score of training set:0.54

```
In [ ]: from sklearn.metrics import r2_score
        p = r2_score(y_test, y_test_pred)
        print('R2 Score of testing set:%.2f'%p)
```

R2 Score of testing set:0.51